

In [1]:

```
1 import os
2 import tensorflow as tf
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import time
6
7 %matplotlib inline
```

In [2]:

```
1 tf.compat.v1.logging.set_verbosity(tf.compat.v1.logging.ERROR)
```

In []:

```
1 ##### Data #####
2
3 BATCH_SIZE = 64
4 DATASET = "CIFAR100"
5 NUM_CLASSES = 100
6
7 def normalize_data(X_train, X_test):
8     _mean = X_train.mean(axis=(0, 1, 2), keepdims=True)
9     _std = X_train.std(axis=(0, 1, 2), keepdims=True)
10    X_train = (X_train - _mean) / _std
11    X_test = (X_test - _mean) / _std
12    return X_train, X_test
13
14 def get_tf_dataset_iter(data):
15     dataset = tf.data.Dataset.from_tensor_slices(data)
16     dataset = dataset.repeat(None) # Repeat Indefinitely
17     dataset = dataset.batch(BATCH_SIZE)
18     dataset = dataset.prefetch(1)
19     dataiter = iter(dataset)
20    return dataiter
```

In [3]:

```
1 def get_data(dataset):
2     if dataset == "CIFAR100":
3         data = tf.keras.datasets.cifar100.load_data()
4     elif dataset == "CIFAR10":
5         data = tf.keras.datasets.cifar10.load_data()
6     else:
7         print ("Invalid Dataset!!")
8         exit(-1)
9
10    (X_train, y_train), (X_test, y_test) = data
11    X_train = np.asarray(X_train, dtype=np.float32)
12    y_train = np.asarray(y_train, dtype=np.int32).flatten()
13    X_test = np.asarray(X_test, dtype=np.float32)
14    y_test = np.asarray(y_test, dtype=np.int32).flatten()
15
16    X_train, X_test = normalize_data(X_train, X_test)
17
18    X_train, X_val = np.split(X_train, [int(0.98*len(X_train))])
19    y_train, y_val = np.split(y_train, [int(0.98*len(y_train))])
20
21    return X_train, y_train, X_val, y_val, X_test, y_test
22
23
24 X_train, y_train, X_val, y_val, X_test, y_test = get_data(DATASET)
25 X_train_iter = get_tf_dataset_iter(X_train)
26 y_train_iter = get_tf_dataset_iter(y_train)
27 X_val_iter = get_tf_dataset_iter(X_val)
28 y_val_iter = get_tf_dataset_iter(y_val)
29 X_test_iter = get_tf_dataset_iter(X_test)
30 y_test_iter = get_tf_dataset_iter(y_test)
```

In [4]:

```
1 def timeme(method):
2     def wrapper(*args, **kw):
3         startTime = int(round(time.time() * 1000))
4         result = method(*args, **kw)
5         endTime = int(round(time.time() * 1000))
6
7         print("Execution Time:", endTime - startTime, 'ms')
8         return result
9
10    return wrapper
```

In [5]:

```
1 def get_accuracy(model, Xset, yset, Xiter, yiter):
2     total_batches = (Xset.shape[0] // BATCH_SIZE) + 1
3     num_correct, num_samples = 0, 0
4     for batch_num in range(total_batches):
5         batch_X = next(Xiter)
6         batch_y = next(yiter).numpy()
7         scores = model(batch_X, is_training = False)
8         scores = scores.numpy()
9         pred_y = scores.argmax(axis=1)
10        num_samples += batch_X.shape[0]
11        num_correct += (pred_y == batch_y).sum()
12    accuracy = float(num_correct) / num_samples
13    return 100*accuracy
```

In [31]:

```
1 EPOCH_ITERATIONS = 700
2 LEARNING_RATE = 3e-4
3 EXPERIMENT_ROOT = './experiment'
4 PRINT_FREQUENCY = 100
5
6 @timeme
7 # def train(CNN_class, num_epochs = 10, resume = False, plot_losses = True):
8 def train(model, num_epochs = 10, resume = False, plot_losses = True):
9     # model = CNN_class(NUM_CLASSES)
10     decay_start_iteration = 0.6 * (num_epochs * EPOCH_ITERATIONS)
11     lr_schedule = tf.keras.optimizers.schedules.ExponentialDecay(LEARNING_RATE,
12     optimizer = tf.keras.optimizers.Adam(learning_rate=lr_schedule)
13
14     writer = tf.summary.create_file_writer(EXPERIMENT_ROOT)
15     ckpt = tf.train.Checkpoint(step=tf.Variable(0), optimizer=optimizer, net=model)
16     manager = tf.train.CheckpointManager(ckpt, EXPERIMENT_ROOT, max_to_keep=10)
17
18     if resume:
19         ckpt.restore(manager.latest_checkpoint)
20
21     stop_training = False
22     for epc in range(1, num_epochs + 1):
23         print ("[Epoch {}]".format(epc))
24         loss_list = []
25         validation accuracies = []
26         for iter_num in range(EPOCH_ITERATIONS):
27             batch_X = next(X_train_iter)
28             batch_y = next(y_train_iter)
29             with tf.GradientTape() as tape:
30                 scores = model(batch_X, is_training = True)
31                 losses = tf.nn.sparse_softmax_cross_entropy_with_logits(labels=batch_y, logits=scores)
32                 meanloss = tf.reduce_mean(losses)
33                 loss_list.append(meanloss)
34                 lossnp = losses.numpy()
35
36             if iter_num % PRINT_FREQUENCY == 0:
37
38                 with writer.as_default():
39                     tf.summary.scalar("loss", meanloss, step=iter_num)
40                     tf.summary.histogram('losses', lossnp, step=iter_num)
41
42                 print('iter:{:6d}, loss min|avg|max: {:.3f}|{:.3f}|{:.3f}, '
43                       .format(iter_num,
44                               float(np.min(lossnp)),
45                               float(np.mean(lossnp)),
46                               float(np.max(lossnp))), end="")
47                 validation_acc = get_accuracy(model, X_val, y_val, X_val_iter)
48                 validation accuracies.append(validation_acc)
49                 print ('val acc: {:.2f} %'.format(validation_acc))
50
51                 if epc >= 3 and validation_acc <= 20.0:
52                     print ("Circuit Breaker!! Validation accuracy too low")
53                     stop_training = True
54                     break
55
56         grad = tape.gradient(meanloss, model.trainable_variables)
57         optimizer.apply_gradients(zip(grad, model.trainable_variables))
58         ckpt.step.assign_add(1)
59
```

```

60         manager.save()
61
62         if stop_training:
63             break
64
65         if plot_losses:
66             plt.plot(loss_list, 'r')
67             plt.grid(True)
68             plt.title('Epoch {} Loss'.format(epc))
69             plt.xlabel('MiniBatch number')
70             plt.ylabel('MiniBatch loss')
71             plt.show()
72
73             plt.plot(validation accuracies)
74             plt.grid(True)
75             plt.title('Epoch {} Validation Accuracies'.format(epc))
76             plt.xlabel('MiniBatch number (100s)')
77             plt.ylabel('Validation Accuracy')
78             plt.show()
79     # return model

```

CNN1

In [7]:

```

1  class CNN1(tf.keras.Model):
2      def __init__(self, num_classes, input_shape = (32, 32, 3)):
3          super().__init__()
4          self.conv1 = tf.keras.layers.Conv2D(input_shape = input_shape, filters
5              padding = 'SAME', activation = tf.nn.relu, name = '
6          self.conv2 = tf.keras.layers.Conv2D(filters = 32, kernel_size = [5,5],
7              padding = 'SAME', activation = tf.nn.relu, name = '
8          self.flatten1 = tf.keras.layers.Flatten()
9          self.fc1 = tf.keras.layers.Dense(512, name = 'fc1')
10         self.fc2 = tf.keras.layers.Dense(num_classes, name = 'fc2')
11
12         def call(self, batch_input, is_training):
13             x = self.conv1(batch_input)
14             x = self.conv2(x)
15             x = self.flatten1(x)
16             x = self.fc1(x)
17             x = self.fc2(x)
18         return x

```

CNN2

In [8]:

```
1 class CNN2(tf.keras.Model):
2     def __init__(self, num_classes, input_shape = (32, 32, 3)):
3         super().__init__()
4         self.conv1 = tf.keras.layers.Conv2D(input_shape = input_shape, filters
5             padding = 'SAME', activation = tf.nn.relu, name = '
6         self.conv2 = tf.keras.layers.Conv2D(filters = 32, kernel_size = [3,3],
7             padding = 'SAME', activation = tf.nn.relu, name = '
8         self.flatten1 = tf.keras.layers.Flatten()
9         self.fc1 = tf.keras.layers.Dense(512, name = 'fc1')
10        self.fc2 = tf.keras.layers.Dense(num_classes, name = 'fc2')
11
12    def call(self, batch_input, is_training):
13        x = self.conv1(batch_input)
14        x = self.conv2(x)
15        x = self.flatten1(x)
16        x = self.fc1(x)
17        x = self.fc2(x)
18        return x
```

CNN3

In [9]:

```
1 class CNN3(tf.keras.Model):
2     def __init__(self, num_classes, input_shape = (32, 32, 3)):
3         super().__init__()
4         self.conv1 = tf.keras.layers.Conv2D(input_shape = input_shape, filters
5             padding = 'SAME', activation = tf.nn.relu, name = '
6         self.conv2 = tf.keras.layers.Conv2D(filters = 64, kernel_size = [5,5],
7             padding = 'SAME', activation = tf.nn.relu, name = '
8         self.flatten1 = tf.keras.layers.Flatten()
9         self.fc1 = tf.keras.layers.Dense(512, name = 'fc1')
10        self.fc2 = tf.keras.layers.Dense(num_classes, name = 'fc2')
11
12    def call(self, batch_input, is_training):
13        x = self.conv1(batch_input)
14        x = self.conv2(x)
15        x = self.flatten1(x)
16        x = self.fc1(x)
17        x = self.fc2(x)
18        return x
```

CNN4

In [10]:

```
1 class CNN4(tf.keras.Model):
2     def __init__(self, num_classes, input_shape = (32, 32, 3)):
3         super().__init__()
4         self.conv1 = tf.keras.layers.Conv2D(input_shape = input_shape, filters
5             padding = 'SAME', activation = tf.nn.relu, name = '
6         self.max_pool1 = tf.keras.layers.MaxPooling2D(pool_size = [2,2], stride
7         self.conv2 = tf.keras.layers.Conv2D(filters = 64, kernel_size = [3,3],
8             padding = 'SAME', activation = tf.nn.relu, name = '
9         self.max_pool2 = tf.keras.layers.MaxPooling2D(pool_size = [2,2], stride
10        self.flatten1 = tf.keras.layers.Flatten()
11        self.fc1 = tf.keras.layers.Dense(512, name = 'fc1')
12        self.bnorm1 = tf.keras.layers.BatchNormalization(axis=1,momentum=0.9,epsilon=0.001)
13        self.fc2 = tf.keras.layers.Dense(num_classes, name = 'fc2')
14
15    def call(self, batch_input, is_training):
16        x = self.conv1(batch_input)
17        x = self.max_pool1(x)
18        x = self.conv2(x)
19        x = self.max_pool2(x)
20        x = self.flatten1(x)
21        x = self.fc1(x)
22        x = self.bnorm1(x)
23        x = self.fc2(x)
24        return x
```

CNN5

In [11]:

```
1 class CNN5(tf.keras.Model):
2     def __init__(self, num_classes, input_shape = (32, 32, 3)):
3         super().__init__()
4         self.conv1 = tf.keras.layers.Conv2D(input_shape = input_shape, filters
5             padding = 'SAME', activation = tf.nn.relu, name = '
6         self.max_pool1 = tf.keras.layers.MaxPooling2D(pool_size = [2,2], stride
7         self.conv2 = tf.keras.layers.Conv2D(filters = 64, kernel_size = [3,3],
8             padding = 'SAME', activation = tf.nn.relu, name = '
9         self.max_pool2 = tf.keras.layers.MaxPooling2D(pool_size = [2,2], stride
10        self.flatten1 = tf.keras.layers.Flatten()
11        self.fc1 = tf.keras.layers.Dense(512, name = 'fc1')
12        self.bnorm1 = tf.keras.layers.BatchNormalization(axis=1,momentum=0.9,epsilon=0.001)
13        self.fc2 = tf.keras.layers.Dense(num_classes, name = 'fc2')
14
15    def call(self, batch_input, is_training):
16        x = self.conv1(batch_input)
17        x = self.max_pool1(x)
18        x = self.conv2(x)
19        x = self.max_pool2(x)
20        x = self.flatten1(x)
21        x = self.fc1(x)
22        x = self.bnorm1(x, is_training)
23        x = self.fc2(x)
24        return x
```

CNN6

In [27]:

```
1 class CNN6(tf.keras.Model):
2     def __init__(self, num_classes, input_shape = (32, 32, 3)):
3         super().__init__()
4         self.conv1 = tf.keras.layers.Conv2D(input_shape = input_shape, filters
5             padding = 'SAME', activation = tf.nn.relu, name = '
6         self.bnorm1 = tf.keras.layers.BatchNormalization(axis = 1, momentum = 0
7
8         self.conv2 = tf.keras.layers.Conv2D(filters = 64, kernel_size = [3,3],
9             padding = 'SAME', activation = tf.nn.relu, name = '
10        self.bnorm2 = tf.keras.layers.BatchNormalization(axis = 1, momentum = 0
11
12        self.max_pool1 = tf.keras.layers.MaxPooling2D(pool_size = [2,2], stride
13
14        self.flatten1 = tf.keras.layers.Flatten()
15        self.fc1 = tf.keras.layers.Dense(512, name = 'fc1')
16        self.bnorm3 = tf.keras.layers.BatchNormalization(axis = 1, momentum = 0
17        self.fc2 = tf.keras.layers.Dense(num_classes, name = 'fc2')
18
19    def call(self, batch_input, is_training):
20        x = self.conv1(batch_input)
21        x = self.bnorm1(x, is_training)
22        x = self.conv2(x)
23        x = self.bnorm2(x, is_training)
24        x = self.max_pool1(x)
25        x = self.flatten1(x)
26        x = self.fc1(x)
27        x = self.bnorm3(x, is_training)
28        x = self.fc2(x)
29        return x
```

CNN7

In [13]:

```
1 class CNN7(tf.keras.Model):
2     def __init__(self, num_classes, input_shape = (32, 32, 3)):
3         super().__init__()
4         self.conv1 = tf.keras.layers.Conv2D(input_shape = input_shape, filters
5             padding = 'SAME', activation = tf.nn.relu, name = '
6         self.max_pool1 = tf.keras.layers.MaxPooling2D(pool_size = [2,2], stride
7         self.conv2 = tf.keras.layers.Conv2D(filters = 128, kernel_size = [3,3],
8             padding = 'SAME', activation = tf.nn.relu, name = '
9         self.max_pool2 = tf.keras.layers.MaxPooling2D(pool_size = [2,2], stride
10        self.conv3 = tf.keras.layers.Conv2D(filters = 256, kernel_size = [3,3],
11            padding = 'SAME', activation = tf.nn.relu, name = '
12        self.gpool1 = tf.keras.layers.GlobalAveragePooling2D()
13        self.fc1 = tf.keras.layers.Dense(512, name = 'fc1')
14        self.bnorm1 = tf.keras.layers.BatchNormalization(axis = 1, momentum = 0
15        self.fc2 = tf.keras.layers.Dense(num_classes, name = 'fc2')
16
17    def call(self, batch_input, is_training):
18        x = self.conv1(batch_input)
19        x = self.max_pool1(x)
20        x = self.conv2(x)
21        x = self.max_pool2(x)
22        x = self.conv3(x)
23        x = self.gpool1(x)
24        x = self.fc1(x)
25        x = self.bnorm1(x, is_training)
26        x = self.fc2(x)
27        return x
```

Experiments

Experiment #1 - CNN1

In [46]:

```
1 cnn1_model = CNN1(NUM_CLASSES)
2 train(cnn1_model, num_epochs = 30)
```

[Epoch 1]

```
iter:    0, loss min|avg|max: 4.066|4.752| 5.680, val acc: 1.07 %
iter:   100, loss min|avg|max: 1.429|4.040| 7.318, val acc: 11.23 %
iter:   200, loss min|avg|max: 0.281|3.568| 7.741, val acc: 17.29 %
iter:   300, loss min|avg|max: 0.425|3.223|10.827, val acc: 19.43 %
iter:   400, loss min|avg|max: 0.135|2.976| 6.332, val acc: 24.51 %
iter:   500, loss min|avg|max: 0.030|2.999| 9.676, val acc: 24.80 %
```

KeyboardInterrupt

Traceback (most recent call

last)

<ipython-input-46-458b4dc60723> in <module>

```
1 cnn1_model = CNN1(NUM_CLASSES)
```

```
----> 2 train(cnn1_model, num_epochs = 30)
```

<ipython-input-4-30fa9fee35c5> in wrapper(*args, **kw)

```
2 def wrapper(*args, **kw):
```

```
3     startTime = int(round(time.time() * 1000))
```

```
----> 4     result = method(*args, **kw)
```

```
5     endTime = int(round(time.time() * 1000))
```

```
6
```

<ipython-input-31-4740c8fef31d> in train(model, num_epochs, resume, plot_losses)

```
54         break
```

```
55
```

```
---> 56         grad = tape.gradient(meanloss, model.trainable_variables)
```

```
57         optimizer.apply_gradients(zip(grad, model.trainable_variables))
```

```
58         ckpt.step.assign_add(1)
```

~/local/lib/python3.5/site-packages/tensorflow_core/python/eager/backprop.py in gradient(self, target, sources, output_gradients, unconnected_gradients)

```
1027         output_gradients=output_gradients,
```

```
1028         sources_raw=flat_sources_raw,
```

```
-> 1029         unconnected_gradients=unconnected_gradients)
```

```
1030
```

```
1031     if not self._persistent:
```

~/local/lib/python3.5/site-packages/tensorflow_core/python/eager/imperative_grad.py in imperative_grad(tape, target, sources, output_gradients, sources_raw, unconnected_gradients)

```
75         output_gradients,
```

```
76         sources_raw,
```

```
---> 77         compat.as_str(unconnected_gradients.value))
```

~/local/lib/python3.5/site-packages/tensorflow_core/python/eager/backprop.py in _gradient_function(op_name, attr_tuple, num_inputs, inputs, outputs, out_grads, skip_input_indices)

```
139     return [None] * num_inputs
```

```
140
```

```
--> 141     return grad_fn(mock_op, *out_grads)
```

```
142
```

143

```
~/local/lib/python3.5/site-packages/tensorflow_core/python/ops/math_g
rad.py in _MatMulGrad(op, grad)
    1629     if not t_a and not t_b:
    1630         grad_a = gen_math_ops.mat_mul(grad, b, transpose_b=True)
-> 1631         grad_b = gen_math_ops.mat_mul(a, grad, transpose_a=True)
    1632     elif not t_a and t_b:
    1633         grad_a = gen_math_ops.mat_mul(grad, b)

~/local/lib/python3.5/site-packages/tensorflow_core/python/ops/gen_ma
th_ops.py in mat_mul(a, b, transpose_a, transpose_b, name)
    5604     _ctx._context_handle, tld.device_name, "MatMul", name,
    5605     tld.op_callbacks, a, b, "transpose_a", transpose_a, "t
ranspose_b",
-> 5606     transpose_b)
    5607     return _result
    5608     except _core._FallbackException:
```

KeyboardInterrupt:

In [47]:

```
1 test_acc1 = get_accuracy(cnn1_model, X_test, y_test, X_test_iter, y_test_iter)
2 print ("CNN1 - Test Accuracy: {:.2f} %".format(test_acc2), end="\n\n")
3
4 cnn1_model.summary()
```

CNN1 - Test Accuracy: 30.16 %


Model: "cn_n1"

Layer (type)	Output Shape	Param #
=====		
conv1 (Conv2D)	multiple	1216
conv2 (Conv2D)	multiple	12832
flatten_13 (Flatten)	multiple	0
fc1 (Dense)	multiple	16777728
fc2 (Dense)	multiple	51300
=====		
Total params: 16,843,076		
Trainable params: 16,843,076		
Non-trainable params: 0		

Experiment #2 - CNN2

In [44]:

```
1 cnn2_model = CNN2(NUM_CLASSES)
2 train(cnn2_model, num_epochs = 30)
```



[Epoch 2]
iter: 0, loss min|avg|max: 0.007|2.647| 6.875, val acc: 29.79 %
iter: 100, loss min|avg|max: 0.043|2.798| 7.152, val acc: 29.88 %
iter: 200, loss min|avg|max: 0.004|2.500|12.414, val acc: 31.74 %
iter: 300, loss min|avg|max: 0.000|2.217| 6.415, val acc: 30.76 %
iter: 400, loss min|avg|max: 0.004|1.897| 7.133, val acc: 29.59 %

In [45]:

```
1 test_acc2 = get_accuracy(cnn2_model, X_test, y_test, X_test_iter, y_test_iter)
2 print ("CNN2 - Test Accuracy: {:.2f} %".format(test_acc2), end="\n\n")
3
4 cnn2_model.summary()
```

CNN2 - Test Accuracy: 30.16 %

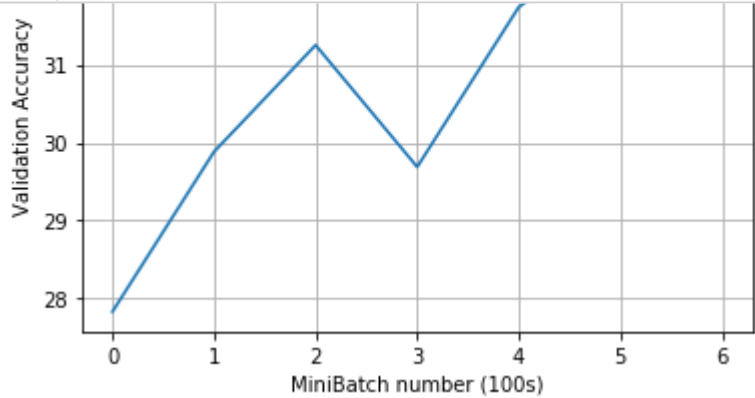
Model: "cn_n2"

Layer (type)	Output Shape	Param #
=====		
conv1 (Conv2D)	multiple	448
conv2 (Conv2D)	multiple	4640
flatten_12 (Flatten)	multiple	0
fc1 (Dense)	multiple	16777728
fc2 (Dense)	multiple	51300
=====		
Total params: 16,834,116		
Trainable params: 16,834,116		
Non-trainable params: 0		

Experiment #3 - CNN3

In [42]:

```
1 cnn3_model = CNN3(NUM_CLASSES)
2 train(cnn3_model, num_epochs = 30)
```



[Epoch 3]
iter: 0, loss min|avg|max: 0.007|1.983|10.436, val acc: 33.59 %

In [43]:

```
1 test_acc3 = get_accuracy(cnn3_model, X_test, y_test, X_test_iter, y_test_iter)
2 print ("CNN3 - Test Accuracy: {:.2f} %".format(test_acc3), end="\n\n")
3
4 cnn3_model.summary()
```

CNN3 - Test Accuracy: 32.90 %


Model: "cn_n3_1"

Layer (type)	Output Shape	Param #
=====		
conv1 (Conv2D)	multiple	2432
conv2 (Conv2D)	multiple	51264
flatten_11 (Flatten)	multiple	0
fc1 (Dense)	multiple	33554944
fc2 (Dense)	multiple	51300
=====		
Total params: 33,659,940		
Trainable params: 33,659,940		
Non-trainable params: 0		

Experiment #4 - CNN4

In [40]:

```
1 cnn4_model = CNN4(NUM_CLASSES)
2 train(cnn4_model, num_epochs = 30)
```



In [41]:

```
1 test_acc4 = get_accuracy(cnn4_model, X_test, y_test, X_test_iter, y_test_iter)
2 print ("CNN4 - Test Accuracy: {:.2f} %".format(test_acc4), end="\n\n")
3
4 cnn4_model.summary()
```

CNN4 - Test Accuracy: 37.42 %


Model: "cn_n4_4"

Layer (type)	Output Shape	Param #
=====		
conv1 (Conv2D)	multiple	896
max_pool1 (MaxPooling2D)	multiple	0
conv2 (Conv2D)	multiple	18496
max_pool2 (MaxPooling2D)	multiple	0
flatten_10 (Flatten)	multiple	0
fc1 (Dense)	multiple	2097664
bnorm1 (BatchNormalization)	multiple	2048
fc2 (Dense)	multiple	51300
=====		
Total params: 2,170,404		
Trainable params: 2,169,380		
Non-trainable params: 1,024		

Experiment #5 - CNN5

In [38]:

```
1 cnn5_model = CNN5(NUM_CLASSES)
2 train(cnn5_model, num_epochs = 30)
```



[Epoch 7]

```
iter: 0, loss min|avg|max: 0.001|1.543| 8.336, val acc: 36.82 %
iter: 100, loss min|avg|max: 0.005|1.142| 5.760, val acc: 35.04 %
```

In [39]:

```
1 test_acc5 = get_accuracy(cnn5_model, X_test, y_test, X_test_iter, y_test_iter)
2 print ("CNN5 - Test Accuracy: {:.2f} %".format(test_acc5), end="\n\n")
3
4 cnn5_model.summary()
```

CNN5 - Test Accuracy: 39.20 %


Model: "cn_n5_1"

Layer (type)	Output Shape	Param #
=====		
conv1 (Conv2D)	multiple	896
max_pool1 (MaxPooling2D)	multiple	0
conv2 (Conv2D)	multiple	18496
max_pool2 (MaxPooling2D)	multiple	0
flatten_9 (Flatten)	multiple	0
fc1 (Dense)	multiple	2097664
bnorm1 (BatchNormalization)	multiple	2048
fc2 (Dense)	multiple	51300
=====		
Total params: 2,170,404		
Trainable params: 2,169,380		
Non-trainable params: 1,024		

Experiment #6 - CNN6

In [32]:

```
1 cnn6_model = CNN6(NUM_CLASSES)
2 train(cnn6_model, num_epochs = 30)
```



[Epoch 5]

iter: 0, loss min|avg|max: 0.001|0.770| 6.122, val acc: 36.04 %

In [34]:

```
1 test_acc6 = get_accuracy(cnn6_model, X_test, y_test, X_test_iter, y_test_iter)
2 print ("CNN6 - Test Accuracy: {:.2f} %".format(test_acc6), end="\n\n")
3
4 cnn6_model.summary()
```

CNN6 - Test Accuracy: 35.65 %

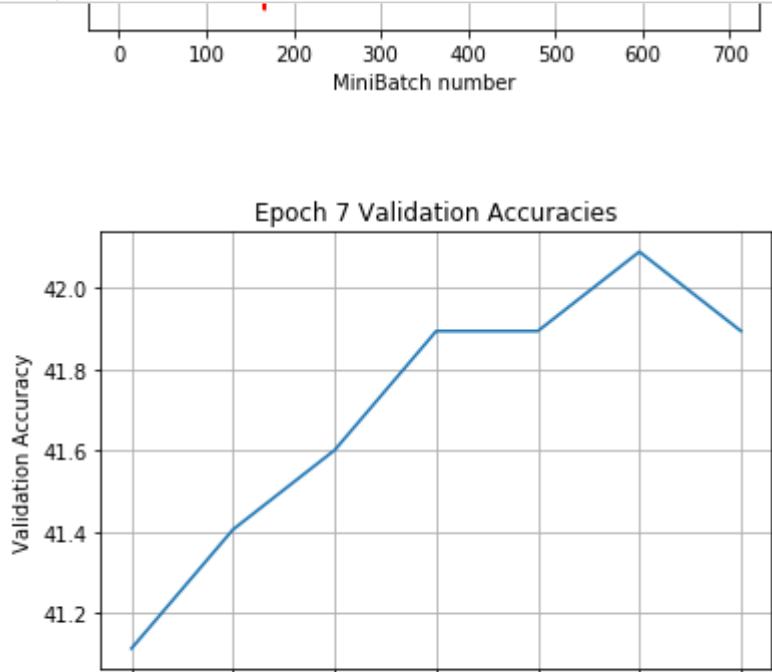
Model: "cn_n6_3"

Layer (type)	Output Shape	Param #
=====		
conv1 (Conv2D)	multiple	896
bnorm1 (BatchNormalization)	multiple	128
conv2 (Conv2D)	multiple	18496
bnorm2 (BatchNormalization)	multiple	128
max_pool1 (MaxPooling2D)	multiple	0
flatten_8 (Flatten)	multiple	0
fc1 (Dense)	multiple	8389120
bnorm3 (BatchNormalization)	multiple	2048
fc2 (Dense)	multiple	51300
=====		
Total params: 8,462,116		
Trainable params: 8,460,964		
Non-trainable params: 1,152		
=====		

Experiment #7 - CNN7

In [35]:

```
1 cnn7_model = CNN7(NUM_CLASSES)
2 train(cnn7_model, num_epochs = 30)
```



In [36]:

```
1 test_acc7 = get_accuracy(cnn7_model, X_test, y_test, X_test_iter, y_test_iter)
2 print ("CNN7 - Test Accuracy: {:.2f} %".format(test_acc7), end="\n\n")
3
4 cnn7_model.summary()
```

CNN7 - Test Accuracy: 41.67 %

Model: "cn_n7"

Layer (type)	Output Shape	Param #
=====		
conv1 (Conv2D)	multiple	1792
max_pool1 (MaxPooling2D)	multiple	0
conv2 (Conv2D)	multiple	73856
max_pool2 (MaxPooling2D)	multiple	0
conv3 (Conv2D)	multiple	295168
global_average_pooling2d_1 (multiple	0
fc1 (Dense)	multiple	131584
bnorm1 (BatchNormalization)	multiple	2048
fc2 (Dense)	multiple	51300
=====		
Total params: 555,748		
Trainable params: 554,724		
Non-trainable params: 1,024		

Max. Test Accuracy = 41.67 %

CNN architecture - **CNN7**

Epochs: **7**

Iterations per epoch: **700**

Learning Rate: **3e-4**